



# TrustLens

AI-Powered Community Trust & Investigation Platform

*"Don't trust the name. Investigate the evidence."*

**Team: POCKETMIND**

iQOO Hackathon 2026 — Grand Finale

TARGET TRACK

**Community App**

Project Document · Current Prototype + Grand Finale Build Plan

A general-purpose community trust concept, demonstrated on an iQOO phone.

# Contents

01 Executive Summary . . . . .

02 Problem Statement . . . . .

03 Existing Approaches . . . . .

04 The Gap . . . . .

05 TrustLens Solution . . . . .

06 What Makes TrustLens Unique . . . . .

07 Community App Track Fit . . . . .

08 Community Intelligence . . . . .

09 AI at the Core . . . . .

10 On-Device AI Direction . . . . .

11 Phone-First Experience . . . . .

12 Technical Architecture . . . . .

13 Risk Engine . . . . .

14 Technology Stack . . . . .

15 Security & Privacy . . . . .

16 Current Prototype . . . . .

17 Grand Finale Build Plan . . . . .

18 Demo Scenarios . . . . .

19 Hackathon Evaluation Alignment . . . . .

20 Impact . . . . .

21 Responsible AI . . . . .

22 Limitations . . . . .

23 Future Scope . . . . .

24 Conclusion . . . . .

25 Team POCKETMIND . . . . .

3

3

4

4

5

6

7

8

9

9

10

11

12

12

13

14

14

15

16

16

17

17

17

17

18

## How to read the status labels in this document

|                              |                                                                                                                                                                                                                 |
|------------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| CURRENT PROTOTYPE            | What has already been demonstrated in the existing repository and demo: concept, UI / prototype, investigation workflow, controlled scenarios, risk/result demonstration and the evidence-relationship concept. |
| PLANNED                      | A component or technology chosen for the Grand Finale build but not yet demonstrated. Subject to final implementation decisions.                                                                                |
| GRAND FINALE BUILD DIRECTION | A capability we intend to build during the hackathon (for example the community layer). It is a direction, not a completed feature.                                                                             |

## 01 Executive Summary

**TrustLens** is an AI-powered multimodal community trust and investigation platform. It helps students, developers, professionals, job seekers, event participants and technology communities investigate suspicious digital opportunities, and share useful trust intelligence with each other.

Users provide evidence through the phone they already carry:

- **Camera** – posters, business cards, web pages, payment requests
- **QR codes** – registration and payment links
- **Images, documents and screenshots** – offer letters, PDFs, chats, emails
- **Voice** – natural questions such as “Is this internship genuine?”

TrustLens extracts entities, connects the evidence, verifies relationships, analyses risk and generates an explainable investigation report. Rather than returning a bare verdict, it shows *which* pieces of evidence are verified, mismatched, unknown or risky, and why.

**One person’s investigation can help another person make a safer decision.**

| Item              | Detail                                                                                                                                                                                           |
|-------------------|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Project           | TrustLens — AI-Powered Community Trust & Investigation Platform                                                                                                                                  |
| Team              | POCKETMIND                                                                                                                                                                                       |
| Event / track     | iQOO Hackathon 2026 Grand Finale — Community App                                                                                                                                                 |
| Current status    | <b>CURRENT PROTOTYPE</b> Product concept, UI, investigation workflow and controlled scenarios are demonstrated.                                                                                  |
| Grand Finale goal | <b>GRAND FINALE BUILD DIRECTION</b> A working Android application with camera-based investigation, multimodal AI, verification, risk analysis, an evidence graph and a community evidence layer. |
| Positioning       | A general-purpose community trust tool. It is not specific to iQOO; an iQOO phone is the primary demonstration device.                                                                           |

## 02 Problem Statement

**A friend paid ₹850 for an internship, completed the assigned work, and later discovered that the opportunity was fraudulent.**

This was not a crude scam. Suspicious opportunities increasingly look professional because they borrow real credibility:

- Real company names and real logos
- Fake recruiters, fake websites and fake offer letters
- Payment links and QR codes
- Professional-looking messages and documents

### What the user is forced to verify manually

To make a safe decision the user must independently check the company registration, the website, the email domain, the recruiter, the offer letter, reviews, the payment recipient and the registration link — usually across many separate tools, usually under time pressure.

**A company itself may be legitimate while a specific opportunity using its identity is fake. The problem is not simply company verification — it is opportunity-level trust verification.**

## The community dimension

When one person uncovers useful evidence — a mismatched domain, a suspicious payee, a recruiter who does not exist at the company — that knowledge normally disappears into a private chat or a scattered post. The open question is: **if one person discovers useful evidence, how can that knowledge help others who encounter the same entity or opportunity?**

## 03 Existing Approaches

Several categories of tools address parts of this problem. Each is useful; none is wrong. Their limits come from the narrow unit they inspect.

| Approach                 | What it does                                    | Limitation                                                                                   |
|--------------------------|-------------------------------------------------|----------------------------------------------------------------------------------------------|
| URL / website checkers   | Check domains, URLs or website reputation.      | May not understand the full context of the opportunity, recruiter, offer or payment request. |
| Spam / message detectors | Analyse emails and messages.                    | Generally focus on the message rather than the entire evidence ecosystem.                    |
| Company lookup platforms | Provide registration and business information.  | A legitimate company can still be impersonated or misused.                                   |
| QR scanners              | Decode QR codes and URLs.                       | Decoding a QR code does not prove the complete opportunity is legitimate.                    |
| Generic AI chatbots      | Read text, answer questions, summarise content. | Not necessarily structured as an evidence-relationship and verification system.              |
| Community reports        | Communities share warnings and experiences.     | Reports can be fragmented, unstructured, hard to connect and hard to verify.                 |

## 04 The Gap

Typical tools treat a trust decision as a set of independent lookups: *URL → check, message → check, company → check, QR → check*. Real trust decisions depend on how the pieces relate to each other.




Figure 1 — Isolated checks (left) versus connected evidence (right).

**A real company, a real logo and a real-looking website can all pass individual checks while the recruiter, email domain and payment recipient do not belong to that company. The missing layer is connecting the evidence — and TrustLens is designed to investigate the relationships between these pieces of evidence.**

## 05 TrustLens Solution

An AI-powered community trust and investigation platform that connects multimodal evidence, verifies relationships, analyses risk, and can turn useful investigations into shared community intelligence.

### 5.1 Evidence inputs

| Input  | Examples                                                                                                                             |
|--------|--------------------------------------------------------------------------------------------------------------------------------------|
| Camera | Internship / job posters, event posters, QR codes, business cards, websites, payment requests                                        |
| Upload | Offer letters, PDFs, emails, WhatsApp screenshots, LinkedIn screenshots, event invitations, payment screenshots                      |
| Voice  | "Is this internship genuine?" · "Why is this suspicious?" · "Show me the negative evidence." · "Can I trust this registration link?" |

### 5.2 Core investigation model



Figure 2 — SEE → EXTRACT → IDENTIFY → VERIFY → CROSS-CHECK → ANALYZE → EXPLAIN → ACT.

| Stage       | What happens                                                                          |
|-------------|---------------------------------------------------------------------------------------|
| See         | Capture or upload real-world evidence.                                                |
| Extract     | Use OCR, vision and document processing to pull out text, QR data and structure.      |
| Identify    | Extract companies, recruiters, emails, URLs, opportunities, payments and events.      |
| Verify      | Check available official and external sources.                                        |
| Cross-check | Compare relationships between entities (does this recruiter belong to this company?). |
| Analyze     | Combine verified, mismatched, suspicious and unknown signals.                         |
| Explain     | Generate an understandable, evidence-based report.                                    |
| Act         | Provide recommended next steps before paying, sharing data or registering.            |

## 06 What Makes TrustLens Unique

The differentiation is not “TrustLens uses AI”. It is the combination of what is investigated and how.

### 6.1 Multimodal investigation

Camera, documents, screenshots, QR codes and voice feed one investigation instead of five separate tools. The evidence picture is built from all of them together.

### 6.2 Evidence relationship analysis

TrustLens investigates whether company, recruiter, email, website, offer and payment actually belong together. Each relationship is marked verified, mismatched or unknown.



Figure 3 — Illustrative evidence relationships for a suspicious internship (concept; not a live result).

### 6.3 Opportunity-level verification

A company can be real while an offer using its name is fake. TrustLens therefore asks: **is this particular opportunity actually connected to this company, person or event?**

### 6.4 Community intelligence

One person’s investigation can become structured evidence that helps another user (Section 8). **GRAND FINALE BUILD DIRECTION**

### 6.5 Explainable results

TrustLens does not output a bare “SCAM”. It shows ✓ **Verified**, ✗ **Mismatch**, ? **Unknown** and ⚠ **Risk Signal** — and explains why each was assigned.

### 6.6 Phone-first

The investigation starts where the suspicious message actually arrives: on the user’s phone, through camera, QR and voice.

### 6.7 Typical approaches versus TrustLens

The table compares *typical / fragmented approaches* in general terms. It is not a claim about any specific product.

| Capability                     | Typical / existing approaches | TrustLens (design target) |
|--------------------------------|-------------------------------|---------------------------|
| URL checking                   | ✓                             | ✓                         |
| Message analysis               | ✓                             | ✓                         |
| Company lookup                 | ✓                             | ✓                         |
| QR scanning                    | ✓                             | ✓                         |
| Document understanding         | Separate / limited            | ✓                         |
| Voice interaction              | Limited / separate            | ✓                         |
| Multimodal investigation       | Fragmented                    | ✓                         |
| Cross-evidence relationships   | Usually limited               | Core                      |
| Opportunity-level verification | Limited                       | Core                      |
| Community trust signals        | Fragmented                    | Core direction            |
| Evidence graph                 | Rare                          | Core direction            |
| Explainable investigation      | Varies                        | Core direction            |
| Phone-first investigation      | Varies                        | Core experience           |

The TrustLens column describes the product design target. What is demonstrated today and what is planned is separated in Sections 16 and 17.

## 07 Community App Track Fit

**Track brief: “Build a community app, not specific to iQOO, that connects developers, professionals, or interest groups, with AI (preferred to be on device) at the core.”**

TrustLens is a community app because it connects people who face the same risks — and because trust intelligence becomes more valuable when it is shared.




Figure 4 — TrustLens as a shared trust layer for many communities.

| Track requirement    | How TrustLens responds                                                                                                                 |
|----------------------|----------------------------------------------------------------------------------------------------------------------------------------|
| Community app        | Connects students, developers, professionals, job seekers, event participants and technology communities around shared trust evidence. |
| Not specific to iQOO | General-purpose investigation concept; the iQOO phone is the demonstration device only.                                                |
| AI at the core       | AI performs multimodal extraction, document and relationship understanding, and explanation — not just chat (Section 9).               |
| On-device preferred  | Local AI will be evaluated for suitable tasks such as OCR and pre-processing (Section 10).                                             |

TrustLens is deliberately not a generic social network. Its purpose is sharing and investigating trust-related evidence.

## 08 Community Intelligence

### GRAND FINALE BUILD DIRECTION

The full community layer is a Grand Finale build direction. It is not claimed as complete in the current prototype.

### 8.1 A concrete scenario

**User A** receives a suspicious internship: an offer from a recruiter, pointing to a website, with a payment request. TrustLens investigates, and User A can choose to contribute the useful structured evidence — entities and relationships, not raw personal documents.

**User B** later receives a different offer involving the same recruiter, company or website. TrustLens can match the entities and surface the relevant community signal alongside its own investigation.




Figure 5 — From one investigation to a related signal for the next user.

**Investigate → Contribute → Connect evidence → Help others → Improve community trust intelligence.**

### 8.2 What a community signal would contain **PLANNED**

- Normalised entities (e.g. email domain, website, recruiter name, payment handle, event name)
- Relationship outcomes (verified / mismatch / unknown) and the evidence category behind them
- Risk outcome and date, so stale signals can be discounted
- No raw documents or personal identifiers unless the user explicitly chooses to share them

### 8.3 Safeguards we plan to design in

- **Advisory, not verdicts.** A community signal is context shown to the user; it never overrides the evidence in the current investigation.
- **Consent first.** Contribution is opt-in and the user can review what will be shared.
- **Inaccurate reports are expected.** Signals carry confidence and recency, and can be corroborated, disputed or reported.

## 09 AI at the Core

AI in TrustLens is not a chatbot bolted onto a form. It performs the work that makes relationship-level investigation possible:

- Multimodal understanding of images, documents and voice
- Evidence and entity extraction
- Document understanding (offer letters, payment screenshots)
- Relationship understanding between extracted entities
- Investigation assistance and follow-up question handling
- Explanation generation in plain language
- Community evidence organisation

GRAND FINALE BUILD DIRECTION




Figure 6 — Evidence-to-report pipeline.

### Where AI is used, and where it deliberately is not

| AI-assisted                                                            | Deterministic                                     |
|------------------------------------------------------------------------|---------------------------------------------------|
| Extracting text, entities and claims from images, PDFs and screenshots | Verification rule checks (e.g. domain comparison) |
| Understanding document structure                                       | Risk level calculation                            |
| Suggesting relationships between entities                              | Thresholds and rule weights                       |
| Explaining findings in plain language                                  | Status labels (verified / mismatch / unknown)     |
| Organising evidence for community signals                              | Auditable, repeatable outcomes                    |

**PLANNED** Gemini is the planned multimodal component for understanding and explanation. No component is claimed to be production-ready; reliability will be established through controlled scenarios during the build.

## 10 On-Device AI Direction

GRAND FINALE BUILD DIRECTION

The Community App track prefers AI that runs on the device. The Grand Finale build will evaluate local AI for suitable tasks and keep cloud AI for tasks that need greater reasoning capacity.




Figure 7 — Hybrid local / cloud investigation architecture.

| Candidate task             | Why local is attractive                                    | Evaluation question                            |
|----------------------------|------------------------------------------------------------|------------------------------------------------|
| OCR (ML Kit)               | Text can be read without sending the image off the device. | Accuracy on posters, letters, low-light photos |
| QR / barcode detection     | Fast and fully local.                                      | Reliability on printed and on-screen codes     |
| Initial image processing   | Cropping, quality checks and redaction before upload.      | Latency and battery impact                     |
| Lightweight classification | Quick triage of evidence type.                             | Accuracy against cloud model                   |
| Local text processing      | Normalise entities before any network call.                | Coverage of messy real-world text              |
| Voice pre-processing       | Reduces data sent and speeds responses.                    | Language / accent coverage                     |

**On-device AI will be evaluated and integrated where technically practical during the Grand Finale. TrustLens does not claim that all AI currently runs on the device.**

## 11 Phone-First Experience

Suspicious opportunities arrive on phones and are acted on from phones. The phone is therefore not just a display — it is the investigation instrument.




Figure 8 — The phone as capture device, question interface and report surface.

| Phone capability                   | Role in TrustLens                                                                  |
|------------------------------------|------------------------------------------------------------------------------------|
| Camera                             | Capture posters, documents, business cards and web pages in the real world.        |
| QR scanning                        | Inspect registration and payment QR codes before any money moves.                  |
| Voice                              | Ask and follow up hands-free: "Show me the negative evidence."                     |
| Screenshots / files <b>PLANNED</b> | Bring in chats, emails and offer letters already on the phone.                     |
| Phone ↔ laptop <b>PLANNED</b>      | Review longer reports or the evidence graph on a larger screen (Office Kit usage). |

## 12 Technical Architecture

The Android application handles capture, interaction and suitable local processing. A FastAPI backend orchestrates AI, verification sources, community intelligence, the investigation and risk engines, and the database.




Figure 9 — Layered architecture (planned Grand Finale build).

### 12.1 Modules

| Module                 | Purpose                                            |
|------------------------|----------------------------------------------------|
| User & Investigation   | Starts and manages investigations                  |
| Camera & Upload        | Captures evidence                                  |
| OCR & QR               | Extracts text and QR data                          |
| Entity Extraction      | Identifies companies, recruiters, offers, payments |
| Web Investigation      | Retrieves external evidence                        |
| Verification           | Checks relationships and evidence                  |
| Community Intelligence | Connects related trust signals                     |
| Cross-Check            | Finds inconsistencies                              |
| Risk Engine            | Calculates consistent risk                         |
| AI Explanation         | Explains findings                                  |
| Trust Report           | Presents the final investigation                   |

## 12.2 Entity resolution and evidence graph **PLANNED**

Extracted entities are normalised (for example, email → domain, URL → registrable domain, payment handle → payee) and matched against each other and against verified sources. Entities become nodes; relationships become edges labelled verified, mismatch or unknown. The same graph supports the explainable report and community matching.

## 12.3 Backend and data **PLANNED** subject to final implementation decision

- **Backend:** Python + FastAPI as the API layer between the Android app and investigation services
- **Database:** PostgreSQL via Supabase, as in the earlier design, with authentication and storage
- **Stored data:** investigation records, extracted entities, evidence metadata, risk results and investigation history
- **Storage:** secure object storage for uploaded evidence, with limited retention (Section 15)

# 13 Risk Engine

**Risk should not depend only on an AI opinion. TrustLens uses a deterministic, rule-based layer so the same evidence always produces the same risk level.**

AI extracts and explains; rules decide. Each signal is classified (verified, mismatch, unknown or risk signal), the rule set combines them, and the result maps to one of three levels. Rule weights and thresholds **PLANNED** will be tuned and tested against controlled scenarios during the build.

| Risk level         | Meaning                                                     |
|--------------------|-------------------------------------------------------------|
| <b>LOW RISK</b>    | Evidence is broadly consistent.                             |
| <b>MEDIUM RISK</b> | Some information needs additional verification.             |
| <b>HIGH RISK</b>   | Multiple significant inconsistencies or suspicious signals. |

## Illustrative signals

| Signal                                                | Typical classification |
|-------------------------------------------------------|------------------------|
| Company identity found in an official source          | ✓ <b>Verified</b>      |
| Recruiter not linked to the company                   | ✗ <b>Mismatch</b>      |
| Email domain differs from the official company domain | ✗ <b>Mismatch</b>      |
| Upfront registration or training fee requested        | ⚠ <b>Risk signal</b>   |
| Payment recipient relationship not found              | ? <b>Unknown</b>       |

**High Risk does not mean proven fraud. It is an evidence-based assessment that tells the user to verify through official channels before acting.**

# 14 Technology Stack

The stack below is the **PLANNED** Grand Finale stack, carried over from the earlier design. The prototype demonstrates the concept and workflow; each layer will be confirmed as it is implemented.

| Layer         | Technology                                      | Purpose                       |
|---------------|-------------------------------------------------|-------------------------------|
| Mobile        | Android + Kotlin + Jetpack Compose              | Phone-native application      |
| Camera        | CameraX                                         | Camera integration            |
| OCR           | Google ML Kit                                   | Local text recognition        |
| QR            | ML Kit Barcode Scanning                         | QR / payment QR detection     |
| Voice         | Android SpeechRecognizer                        | Voice input                   |
| AI            | Gemini API / suitable local model               | Multimodal understanding      |
| Backend       | Python + FastAPI                                | API layer                     |
| Database      | Supabase PostgreSQL (subject to final decision) | Data, authentication, storage |
| Investigation | Search / website / domain APIs                  | External evidence             |
| Risk          | Python deterministic engine                     | Consistent risk calculation   |
| Explanation   | Gemini / suitable model                         | Explain findings              |
| Deployment    | Render / Railway                                | FastAPI hosting               |
| File storage  | Supabase Storage                                | Uploaded evidence             |

## 15 Security & Privacy

These are the design principles for the Grand Finale build. Evidence is personal and sometimes sensitive, so privacy is a design constraint, not an add-on.

| Principle                  | Approach                                                                   |
|----------------------------|----------------------------------------------------------------------------|
| On-device processing       | Process locally where practical to improve privacy and responsiveness.     |
| Data minimisation          | Send only what an investigation needs to the backend or external services. |
| Secure communication       | Encrypted API traffic between app and backend.                             |
| Authenticated access       | Users can reach only their own investigations and evidence.                |
| Secure storage             | Uploaded evidence and investigation data stored with access controls.      |
| Protected credentials      | API keys and service credentials kept server-side, never in the app.       |
| Controlled file processing | Uploaded files validated and handled in a controlled way.                  |
| Sensitive logging controls | Avoid logging personal evidence content.                                   |
| Limited retention          | Evidence kept only as long as needed.                                      |
| User control               | Users decide what is stored, deleted or contributed to the community.      |

**Unknown information should not automatically be treated as fraud.**

## 16 Current Prototype

### CURRENT PROTOTYPE

The current repository is a demo / prototype, not the completed product. This section states plainly what exists.

| Capability                         | Status in current demo |
|------------------------------------|------------------------|
| Product concept                    | ✓ Demonstrated         |
| UI / prototype                     | ✓ Demonstrated         |
| Investigation workflow             | ✓ Demonstrated         |
| Controlled investigation scenarios | ✓ Demonstrated         |
| Risk / result demonstration        | ✓ Demonstrated         |
| Evidence relationship concept      | ✓ Demonstrated         |
| Phone-first interaction concept    | ✓ Demonstrated         |

### What the demo proves

- The investigation workflow is understandable on a phone
- Showing verified, mismatched and unknown evidence is clearer than a single verdict
- Relationship-level reasoning is a viable way to present a trust investigation

### What the demo does not claim

| Not claimed as complete                    |
|--------------------------------------------|
| Full community backend                     |
| Complete verification infrastructure       |
| Full production risk engine                |
| Complete on-device AI                      |
| Complete real-world investigation platform |

## 17 Grand Finale Build Plan

### GRAND FINALE BUILD DIRECTION

Planned deliverables for the hackathon:

| Area          | Planned build                                                           |
|---------------|-------------------------------------------------------------------------|
| Application   | Real Android application (Kotlin, Jetpack Compose)                      |
| Capture       | Camera-based investigation; QR scanning; screenshot and document upload |
| AI            | Multimodal AI, OCR / document understanding, voice interaction          |
| Investigation | Entity resolution, evidence verification, investigation engine          |
| Reasoning     | Dynamic risk analysis, evidence graph, explainable trust report         |
| Community     | Community layer and community evidence signals                          |
| Evaluation    | Evaluation of on-device / local AI; strong phone-first experience       |




Figure 10 — Eight-phase build sequence.

| Phase                     | Focus                                                        | Outcome                                         |
|---------------------------|--------------------------------------------------------------|-------------------------------------------------|
| 1 Core Investigation      | Capture / upload → OCR / QR → entity extraction on Android   | Working end-to-end skeleton on a real phone     |
| 2 Verification            | Official-source and website / domain lookups                 | Verified / mismatch / unknown per relationship  |
| 3 Risk Engine             | Deterministic rules, thresholds, tests                       | Consistent, repeatable risk levels              |
| 4 Evidence Graph          | Entity resolution and relationship edges                     | Graph view that explains the result             |
| 5 Community Layer         | Structured signals, opt-in contribution, entity matching     | Related community signal shown to a second user |
| 6 Phone Features          | Voice, share-in of files and screenshots, camera / QR polish | Natural phone-first interaction                 |
| 7 On-Device AI Evaluation | Benchmark local vs cloud for suitable tasks                  | Local AI integrated where practical             |
| 8 Final UX + Demo         | Polish, rehearsal, fallback demo path                        | Reliable live demo                              |

**A reliable end-to-end workflow is more important than adding many unfinished features.**

## 18 Demo Scenarios

### 18.1 Event verification — the iQOO Hackathon poster

The user points the camera at an event poster and asks whether the registration link can be trusted.




Figure 11 — Poster investigation flow.

### 18.2 Broader Grand Finale scenario — suspicious internship GRAND FINALE BUILD DIRECTION




Figure 12 — Illustrative internship investigation outcome (example; not a live result).

**HIGH RISK ≠ PROVEN FRAUD.** The report tells the user which evidence to verify through official channels before paying or sharing sensitive information.

## 19 Hackathon Evaluation Alignment

| Criterion           | Weight      | TrustLens alignment                                                                             |
|---------------------|-------------|-------------------------------------------------------------------------------------------------|
| End Product Quality | 30%         | A complete community → investigation → result workflow that works reliably on a real phone.     |
| Novelty & Impact    | 20%         | Community-powered trust intelligence built on relationship-level evidence, not isolated checks. |
| Creative Phone Use  | 15%         | Camera + QR + voice and phone-first investigation at the moment a suspicious offer arrives.     |
| Technical Depth     | 15%         | Multimodal AI, verification, entity resolution, deterministic risk engine and evidence graph.   |
| Office Kit Usage    | 10%         | Phone ↔ laptop workflow for capture and review.                                                 |
| Demo & Presentation | 10%         | A real-world story — the ₹850 internship — told end to end.                                     |
| <b>Total</b>        | <b>100%</b> |                                                                                                 |

## 20 Impact

| Who                       | How TrustLens helps                                                        |
|---------------------------|----------------------------------------------------------------------------|
| <b>Students</b>           | Investigate internships and job offers before paying or sharing documents. |
| <b>Developers</b>         | Verify developer events, communities and opportunities.                    |
| <b>Professionals</b>      | Investigate unfamiliar business or professional requests.                  |
| <b>Job seekers</b>        | Investigate recruiters and offers.                                         |
| <b>Event participants</b> | Verify organisers and registration links.                                  |
| <b>Communities</b>        | Share useful trust evidence in a structured, reusable form.                |

**Turn individual investigations into collective trust intelligence.**

## 21 Responsible AI

- **Explainable results** – every finding is tied to evidence
- **Evidence-backed reasoning** – conclusions are not asserted without support
- **Clear uncertainty** – unknown is shown as unknown
- **Human decision-making** – TrustLens informs; the user decides
- **No automatic accusation** – entities are never labelled criminal
- **No guaranteed fraud claims** – the system never promises certainty

The report distinguishes five states instead of treating every unknown signal as fraud:

| Verified                       | Unknown                | Inconsistent       | Suspicious               | High Risk                    |
|--------------------------------|------------------------|--------------------|--------------------------|------------------------------|
| Confirmed by a reliable source | Not enough information | Evidence conflicts | Pattern warrants caution | Multiple significant signals |

## 22 Limitations

- External information may be incomplete
- Public information may be outdated
- Some sources may be unavailable
- AI analysis can require additional verification
- Community information may contain inaccurate reports
- Risk assessment cannot guarantee authenticity

**TrustLens provides investigation support, not guaranteed fraud detection.**

## 23 Future Scope

| Investigation                                  | Community and platform           |
|------------------------------------------------|----------------------------------|
| More official data sources                     | Community reputation signals     |
| Advanced website analysis                      | Continuous reputation monitoring |
| Expanded scam-pattern detection                | Evidence graph expansion         |
| Advanced document analysis                     | Multilingual voice               |
| Communication analysis                         | On-device AI improvements        |
| More sophisticated investigation relationships |                                  |

## 24 Conclusion

TrustLens combines:

**Multimodal AI + Phone-first interaction + Evidence verification + Relationship analysis + Deterministic risk assessment + Community intelligence**

The product vision is **a community-powered AI trust layer for the digital world.**




Figure 13 — The TrustLens story.

**Connect the evidence. Understand the relationships. Help the community.**

## 25 Team POCKETMIND

| Team members     |  |
|------------------|--|
| Anjali Neelam    |  |
| Mudigonda Nikhil |  |
| Srujana Sadhu    |  |

### Appendix A — References carried over from the earlier project document

| No. | Reference                                                                                  | Year | Relevance                               |
|-----|--------------------------------------------------------------------------------------------|------|-----------------------------------------|
| 1   | Tahmasebi et al., “Multimodal Misinformation Detection using Large Vision-Language Models” | 2024 | Multimodal evidence & fact verification |
| 2   | Liu et al., “TextMonkey: An OCR-Free Large Multimodal Model for Understanding Document”    | 2024 | Document understanding                  |
| 3   | Greif et al., “Multimodal LLMs for OCR, OCR Post-Correction, and Named Entity Recognition” | 2025 | OCR & entity extraction                 |
| 4   | “Detecting Online Recruitment Fraud Using Machine Learning”                                | 2021 | Recruitment fraud                       |
| 5   | “Fake Job Posting Detection using Machine Learning Algorithms”                             | 2025 | Fake job detection                      |
| 6   | “Phishing Detection: A Recent Intelligent Machine Learning Comparison”                     | 2017 | Phishing & URL analysis                 |
| 7   | Duarte et al., “Machine Learning for Early Detection of Phishing URLs in Parked Domains”   | 2025 | URL / domain verification               |
| 8   | “Evolution of Phishing Detection: From Traditional Methods to Explainable AI”              | 2026 | Explainable AI & risk analysis          |

# TrustLens

AI-Powered Community Trust & Investigation Platform

---

**Connect the evidence.**  
**Understand the relationships.**  
**Help the community.**

**One investigation can protect  
more than one person.**

*Investigate before you trust.*

## Team POCKETMIND

Anjali Neelam  
Mudigonda Nikhil  
Srujana Sadhu